

System Skill Final Quiz

Date: Tuesday, December 7th, 2021

Due: Thursday, December 9th, 2021 at 11.59PM

Instructor: Rachata Ausavarungnirun

Problem 1 (15 Points):	
Problem 2 (20 Points):	
Problem 3 (15 Points):	
Problem 4 (20 Points):	
Problem 5 (30+5 Points):	
Extra Credit (20 Points):	
Total (100+25 Points):	

Instructions:

1. This is a 72-hour exam. **If you get 100, you get a full score. Any points above 100 goes to your extra credit at the conversion rate of 50% per point.**
2. The total points is far greater than 100. This is intended so that 1) you can pick the questions you are more comfortable and 2) allows rooms for extra credit if you know all the class material. Please read all the questions first.
3. Submit your work as a pdf file on Canvas.
4. Clearly indicate your final answer for each conceptual problem.
5. **DO NOT CHEAT.** If we catch you cheating in any shape or form, you will be penalized based on **my plagiarism policy** ($N * 10\%$ of your total grade, where N is the number of times you plagiarized previously).

Tips:

- **Read everything.** Read all the questions on all pages first and formulate a plan.
- **Be cognizant of time.** It is a sad day if you click submit when the submission site close.
- **Canvas allows resubmission.** I will take a look at the last version you submit.
- **Show work when needed.** You will receive partial credit at the instructors' discretion.

Initials: _____

1. Datalab ... Again! (15 points)

We are going to ask you to implement more bit arithmetic manipulation in a similar fashion as assignment 3.

Unlike assignment 3, you are to write the answer on this exam.

Similar to Assignment 3, you are *only* allowed to use the following eight operators:

! ~ & ^ | + << >>

Also, you are not allowed to use any constants longer than 8 bits.

Table 1 describes a set of functions that manipulate and test sets of bits. The “Rating” field gives the difficulty rating (the number of points) for the puzzle, and the “Max ops” field gives the maximum number of operators you are allowed to use to implement each function.

Name	Description	Rating	Max Ops
bitXor(x, y)	return $x \oplus y$ without using only ~ and &	1	14
isEqual(x, y)	return 1 if $x == y$, 0 otherwise	2	5

Table 1: Bit-Level Manipulation Functions.

For every question, **Please explain what you are trying to do.**

Write down the function body for `bitXor(x, y)` below.

```
int bitXor (int x, int y)
{
    int a = x & y;
    int b = ~x & ~y;
    int z = ~a & ~b;
    return z;
}
```

Initials:

Write down the function body for `isEqual(x, y)` below.

```
return !(x^y);
```

For this part, you can only use logical operation and loops. Assume you have a list of pairs in an array `arr[]` of size `n`. However, there is one sad sad sad number in there that does not have it's pair (for example, your array can be `[1, 2, 8, 1, 8]` and the sad number is 2). Write a code to find this sad number. There is no limit on how many operations you want to use but **you can only loop through the entire array once!**. You get zero credit if you loop through the array more than once.

Initials: _____

2. Jump Table [20 points]

In this question, consider the following assembly codes below. We now assume the 32-bit x86 ISA, which upon a function call, the caller will place the first input at `%ebp+8` and the second input at `%ebp+12`. Fill in the rest of the C code for each of the switch cases. Write "NOTHING HERE" if the space should be left blank or if that line of code should not exist (i.e., the program does not suppose to modify `result` at that line).

```
blah:
pushl %ebp
movl %esp, %ebp
movl 8(%ebp), %edx
movl 12(%ebp), %eax
cmpl $7, %edx
ja .L8
jmp *.L9(,%edx,4)
.section .rodata
.align 4
.align 4
.L9:
.long .L8
.long .L4
.long .L5
.long .L5
.long .L8
.long .L7
.long .L6
.long .L4
.text
.L4:
mov (%edx), %eax
jmp .L7
.L5:
mov (%eax), %eax
jmp .L2
.L6:
mov (%edx), %ecx
add %ecx, %eax
jmp .L2
.L7:
addq %eax, %eax
.L8:
incl %eax
.L2:
popl %ebp
ret
```

Initials: _____

In the space below, fill in the blank to reflect the assembly code above.

```
int blah(int a, int b)
{
    int result;
    switch(  a  )
    {
        case   1  :
        case   7  :
            result =   a  ;
            break;
        case   6  :
            result =   b + a  ;
            break;
        case   2  :
        case   3  :
            result =           ;
        case   5  :
            result =   b + b  ;
        default:
            result =   b + b  ;
    }
    return result;
}
```

Initials: _____

3. Code Size [15 points]

For the following code, please fill in the number (in **hexadecimal base**) for the address of each instruction.

Address	Instruction (in binary)	Instruction (in Assembly)
5fa:	55	push %rbp
5fb	48 89 e5	mov %rsp,%rbp
5fe	89 7d fc	mov %edi,-0x4(%rbp)
601	89 75 f8	mov %esi,-0x8(%rbp)
604	8b 55 fc	mov -0x4(%rbp),%edx
607	8b 45 f8	mov -0x8(%rbp),%eax
609	01 d0	add %edx,%eax
60c	5d	pop %rbp
60d	c3	retq
60e:	55	push %rbp
60f	48 89 e5	mov %rsp,%rbp
612	48 83 ec 08	sub \$0x8,%rsp
616	89 7d fc	mov %edi,-0x4(%rbp)
619	89 75 f8	mov %esi,-0x8(%rbp)
61c	8b 55 f8	mov -0x8(%rbp),%edx
61f	8b 45 fc	mov -0x4(%rbp),%eax
632	89 d6	mov %edx,%esi
634	89 c7	mov %eax,%edi
636	e8 cf ff ff ff	callq 5fa

Initials: _____

4. Caching [20 points]

In this question, let's assume that we have a 16-bit system with a single level 2-way set associative cache with 4 sets, and a cache block size of 32 bytes.

How many bits are needed for the setID and the tags? Draw the breakdown of the tag/index/byte-in-block bits.

$$tag = 16 - \log_2(32) - \log_2(4) = 9$$

For the following program, assume that an integer is 4 bytes.

```
int i; // Assume these variables are stored in the registers.
int a[2048]; // Assume that a = 0x1000
int b[2048]; // Assume that b = 0x8000

for(i=0; i<2048; i++)
    a[i] = i;

for(i=0; i<2048; i++)
    b[i] = a[i]++;
```

What is the total number of cache accesses? What is the number of cache hits and what is the number of cache misses? **Show your work.**

$$\begin{aligned} \text{Total number} &= \text{each for loop? } 2048, \text{ it compile 3 time} \\ &= 2048 \times 3 = \underline{6144} \end{aligned}$$

$$\begin{aligned} \text{hits number} &= \text{have to find misses first} = 6144 - 768 \\ &= \underline{5376} \end{aligned}$$

$$\begin{aligned} \text{Misses number} &= \left(\frac{\text{total cache}}{\text{cache block} / \text{each set}} \right) = \frac{6144}{\frac{32}{4}} = \frac{6144}{32} \times 4 \\ &= \underline{768} \end{aligned}$$

Initials: _____

Can I modify this program in order to make sure the program behaves the same but have higher cache hit rate? If so, how? If not, explain why?

Change if to short type

from using 4 bit it will reduce to 2 bit

Probe is $\frac{6144}{16} = 384$ misses

hit: $6144 - 384 = 5760$

more than int type = 384

Initials: _____

5. Virtual Memory [30+5 points]

Let's create a simple **BIG endian** machine that still utilize 2MB page size, and 32-bit address. Assuming the following data in the memory and the page table root is at 0x10, and the page table entries is 32-bit long, where the n most significant bits after the page offset are used for the physical page number.

Address	Values (in hexadecimal) [Lowest bit – Highest bit]
0x00	00 10 20 30 40 50 60 70 80 90 a0 b0 c0 d0 e0 f0
0x10	10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f
0x20	00 10 20 30 40 50 60 70 80 90 a0 b0 c0 d0 e0 f0
0x30	30 31 32 33 34 35 36 37 38 39 3a 3b 3c 3d 3e 3f
0x40	00 10 20 30 40 50 60 70 80 90 a0 b0 c0 d0 e0 f0
0x50	19 15 12 0a 6b 3a 4b 12 91 ac ff fe 3c 3d 3e 4f
0x60	12 50 62 8a 5e 5f df ea 99 ac 74 6b 91 44 33 ef
0x70	70 71 72 73 74 75 76 77 78 79 7a 7b 7c 7d 7e 7f
0x80	91 40 8a 00 8c 14 fe ff 74 13 02 ba 6b 12 4b 31
0x90	90 91 92 93 94 95 96 97 98 99 9a 9b 9c 9d 9e 9f
0xa0	80 00 8a 00 8c 14 fe ff 72 14 0a 6b 10 02 e1 ba
0xb0	30 31 32 33 34 35 36 37 38 39 3a 3b 3c 3d 3e 3f
0xc0	80 00 8a 00 8c 14 fe ff 72 14 0a 6b 10 01 e1 ba
0xd0	91 40 8a 00 8c 14 fe ff 74 13 02 ba 6b 12 4b 31
0xe0	70 00 8a 00 8c 14 fe ff 72 14 0a 6b 10 03 e1 ba
0xf0	91 40 8a 00 8c 14 fe ff 74 13 02 ba 6b 12 4b 31

- (a) What is the physical address for a virtual address 0x0000beef? Put in **Not enough information** if the table does not provide enough information to get the physical address.

$\log_2(2^{20}) = 20$ $\text{Remainder } (32-20) = 12$

Convert to Binary

1011 1110 1116 111

ppn = 0001 0000 0000

So final = 0001 0000 0000 0000 1011 1110 1110 1111

= 0x1000beef

Initials: _____

- (b) What is the physical address for a virtual address $0x0ffffff$? Put in **Not enough information** if the table does not provide enough information to get the physical address.

Not enough information

as 0000 1111 1111 is not in the table

- (c) What is the physical address for a virtual address $0x0000beef$ if this system were to use 16KB page instead? Put in **Not enough information** if the table does not provide enough information to get the physical address.

$$\log_2(16k) = 14$$

$$\text{remaining} = 32 - 14 = 18$$

Beef in 14 bits = 11 1110 1110 1111

PNN = 0001 1000 0001 1001 00

final = 0001 1000 0001 1001 0011 1110 1110 1111

$$8 \times 4 = 32$$

= $0x18193ee f$

Initials:

- (d) Assuming that the memory access takes 100 cycles to access DRAM, the system has 4-level page table (i.e., a page walk have to access the memory 4 times before it can access its data), an TLB access takes 1 cycle, and a L1 cache access to the set takes 1 cycle and the tag comparison in the L1 cache takes another 1 cycle. How long does it takes to load a data that has a TLB miss and a L1 cache hit in a virtualized environment where nested page table is being used? Feel free to explain your answer.

<p>Each cycle</p> $4\text{level} = 100 \times 4 = 400$ $\text{TLB} = 1 \text{ cycle}$ $\text{cache access} = 1 \text{ cycle}$ $t_{\text{tag comparison}} = 1 \text{ cycle}$	<p>cache access & TLB</p> $= 1 \text{ cycle}$ $\therefore 400 + 1 + 1$ $= 402 \text{ cycles}$
---	---

~~20~~

- (e) (Extra credit: 4 points) Assuming the same setup as part d, but in this case, we utilize a container that runs all it's processes under the host machine's user space (i.e., all virtual addresses used in the container are already in the host's user space). What is the latency of a cache access if you incur a TLB miss and an L1 cache hit? Please explain your answer.

$\text{total latency} = [(1\text{level}) (\text{total page walk})] + (\text{TLB} \& t_{\text{tag}})$ $= [(100 \times 2) (5)] + (2 \times 2)$ $= 1004$

Initials:

(f) (Extra credit: 1 points) What are the names of my cats?

“ ໓໓໓ ”

“ ໓໓໓ ”

Initials: _____


6. Extra Credit: 0x44444444 [20 points]

Do not attempt this until you are done with other questions.

A 32-bit processor implements paging-based virtual memory using a single-level page table. The following are the assumptions about the processor's virtual memory.

- A page table entry (PTE) is 4-bytes in size.
- A PTE stores the physical page number in the least-significant bits.
- The base address of the page tables is page-aligned.

The following figure shows the physical memory of the processor at a particular point in time.



4GB Physical Memory

At this point, when the processor executes the following piece of code, it turns out that the processor accesses the page table entry residing at the physical address of 0x44444444.

```
char *ptr = 0x44444444;  
char val = *ptr; // val == 0x44
```

Initials:

What is the page size of the processor? Show work in detail.

Initials:

Log Table

N	$\log_2 N$
1	0
2	1
4	2
8	3
16	4
32	5
64	6
128	7
256	8
512	9
1024 (1k)	10
2048 (2k)	11
4096 (4k)	12
8192 (8k)	13
16384 (16k)	14
32768 (32k)	15
62236 (64k)	16
131072 (128k)	17
262144 (256k)	18
524288 (512k)	19
1048576 (1M)	20
2097152 (2M)	21
4194304 (4M)	22
8388608 (8M)	23
16777216 (16M)	24

Initials:

Stratchpad

Initials:

Stratchpad